

Assignment 10.2

1. Uploaded students.txt using Hadoop

```
[acadgild@localhost sumona-hadoop]$ hadoop fs -put student.txt /user/acadgild/Assignment10
17/11/28 15:44:57 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
[acadgild@localhost sumona-hadoop]$ hadoop fs -ls /user/acadgild/Assignment10
17/11/28 15:45:19 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 1 items
-rw-r--r-- 1 acadgild supergroup 26204 2017-11-28 15:44 /user/acadgild/Assignment10/student.txt
[acadgild@localhost sumona-hadoop]$
```

2. Included few jar files of HBase to the Pig classpath.

```
[acadgild@localhost ~]$ PIG CLASSPATH=/home/acadgild/sumona-hadoop/hbase-0.98.4-hadoop2/lib/hbase-server-0.98.4-hadoop2:/home/acadgild/sumona-hadoop/hbase-0.98.4-hadoop2/lib/hbase-*.jar;
[acadgild@localhost ~]$
```

3. Started the HBase shell and created a table

```
[acadgild@localhost ~]$ start-hbase.sh
starting master, logging to /usr/local/hbase/logs/hbase-acadgild-master-localhost.localdomain.out
[acadgild@localhost ~]$ hbase shell
2017-11-28 16:04:31,727 INFO [main] Configuration.deprecation: hadoop.native.lib is deprecated. Instead, use io.native.lib.available
HBase Shell; enter 'help<RETURN>' for list of supported commands.
Type "exit<RETURN>" to leave the HBase Shell
Version 0.98.14-hadoop2, r4e4aabb93b52f1b0fef6b66edd06ec8923014dec, Tue Aug 25 22:35:44 PDT 2015

hbase(main):001:0> create 'studentAcad','student data'
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/usr/local/hbase/lib/slf4j-log4j12-1.6.4.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/usr/local/hadoop-2.6.0/share/hadoop/common/lib/slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
2017-11-28 16:06:06,216 WARN [main] util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
0 row(s) in 3.5550 seconds

=> Hbase::Table - studentAcad
hbase(main):002:0>
```

4. Load the data from HDFS to Alias relations.

```
grunt> rawD = LOAD '/user/acadgild/Assignment10/student.txt' USING PigStorage()
>> AS (StudentName:chararray,sector:chararray,DOB:chararray,qalification:chararray,score:int,state:chararray, randomName:chararray
);
2017-11-28 16:21:13,838 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - mapred.job.tracker.persist.jobstatus.hours
is deprecated. Instead, use mapreduce.jobtracker.persist.jobstatus.hours
2017-11-28 16:21:13,838 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - mapred.heartbeats.in.second is deprecated.
Instead, use mapreduce.jobtracker.heartbeats.in.second
2017-11-28 16:21:13,838 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - jobclient.completion.poll.interval is depr
ecated. Instead, use mapreduce.client.completion.pollinterval
2017-11-28 16:21:13,839 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - mapred.tasktracker.tasks.sleep-time-before-
sigkill is deprecated. Instead, use mapreduce.tasktracker.tasks.sleep-time-before-sigkill
2017-11-28 16:21:13,839 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - mapred.job.tracker.http.address is depreca
ted. Instead, use mapreduce.jobtracker.http.address
2017-11-28 16:21:13,839 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - dfs.replication.considerLoad is deprecated
. Instead, use dfs.namenode.replication.considerLoad
2017-11-28 16:21:13,839 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - dfs.block.size is deprecated. Instead, use
dfs.blocksize
2017-11-28 16:21:13,839 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - dfs.permissions is deprecated. Instead, us
e dfs.permissions.enabled
2017-11-28 16:21:13,839 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - topology.node.switch.mapping.impl is depre
cated. Instead, use net.topology.node.switch.mapping.impl
2017-11-28 16:21:13,839 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - dfs.access.time.precision is deprecated. I
nstead, use dfs.namenode.access-time.precision
```

5. Now we can transfer the data inside HBase by STORE command.

We need to ensure that we give the correct name for table name created inside HBase. Also the parameters should be kept in mind to avoid mistake.

```
grunt> STORE rawD INTO 'hbase://studentAcad' USING org.apache.pig.backend.hadoop.hbase.HBaseStorage('student data:StudentName,stud
ent_data:sector,student_data:DOB,student_data:qalification,student_data:score,student_data:state,student_data:randomName');
2017-11-28 16:31:52,622 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - mapreduce.job.counters.limit is deprecated
. Instead, use mapreduce.job.counters.max
2017-11-28 16:31:52,622 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - io.bytes.per.checksum is deprecated. Inste
ad, use dfs.bytes-per-checksum
2017-11-28 16:31:52,622 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, us
e fs.defaultFS
2017-11-28 16:31:52,898 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - io.bytes.per.checksum is deprecated. Inste
ad, use dfs.bytes-per-checksum
```

- Once the success message comes as shown below , it is confirmed our data is loaded inside HBase.

```

HadoopVersion  PigVersion      UserId StartedAt      FinishedAt      Features
2.2.0    0.14.0    acadgild      2017-11-28 16:40:31    2017-11-28 16:42:36    UNKNOWN

Success!

Job Stats (time in seconds):
JobId  Maps    Reduces MaxMapTime    MinMapTime    AvgMapTime    MedianMapTime    MaxReduceTime    MinReduceTime    AvgReduceTime
job_1511863758374_0003  1    0    n/a    n/a    n/a    n/a    0    0    0    0    rawD    MAP_ONLY    hb
ase://studentAcad,

Input(s):
Successfully read 0 records from: "/user/acadgild/Assignment10/student.txt"

Output(s):
Successfully stored 0 records in: "hbase://studentAcad"

Counters:
Total records written : 0
Total bytes written : 0
Spillable Memory Manager spill count : 0
Total bags proactively spilled: 0
Total records proactively spilled: 0

```

The result can be displayed through scan command.

```

hbase(main):004:0> scan 'studentAcad'
ROW                                COLUMN+CELL
ABEDNIGO,goverenment,20-10-2000, column=student_data:DOB, timestamp=1511867458220, value=
BBA,100,alabama,madison`
ABEDNIGO,goverenment,20-10-2000, column=student_data:StudentName, timestamp=1511867458220, value=
BBA,100,alabama,madison`
ABEDNIGO,goverenment,20-10-2000, column=student_data:qalification, timestamp=1511867458220, value=
BBA,100,alabama,madison`
ABEDNIGO,goverenment,20-10-2000, column=student_data:score, timestamp=1511867458220, value=
BBA,100,alabama,madison`
ABEDNIGO,goverenment,20-10-2000, column=student_data:sector, timestamp=1511867458220, value=
BBA,100,alabama,madison`
ABEDNIGO,goverenment,20-10-2000, column=student_data:state, timestamp=1511867458220, value=
BBA,100,alabama,madison`
ABROSER,goverenment,18-11-2002,M column=student_data:DOB, timestamp=1511867458026, value=
BBS,3.5,Pennsylvania,prattville*
ABROSER,goverenment,18-11-2002,M column=student_data:StudentName, timestamp=1511867458026, value=
BBS,3.5,Pennsylvania,prattville*
ABROSER,goverenment,18-11-2002,M column=student_data:qalification, timestamp=1511867458026, value=
BBS,3.5,Pennsylvania,prattville*
ABROSER,goverenment,18-11-2002,M column=student_data:score, timestamp=1511867458026, value=
BBS,3.5,Pennsylvania,prattville*
ABROSER,goverenment,18-11-2002,M column=student_data:sector, timestamp=1511867458026, value=
BBS,3.5,Pennsylvania,prattville*
ABROSER,goverenment,18-11-2002,M column=student_data:state, timestamp=1511867458026, value=
BBS,3.5,Pennsylvania,prattville*
AGNES,goverenment,20-10-2000,BE, column=student_data:DOB, timestamp=1511867458221, value=
100,alabama,madison`
AGNES,goverenment,20-10-2000,BE, column=student_data:StudentName, timestamp=1511867458221, value=
100,alabama,madison`
AGNES,goverenment,20-10-2000,BE, column=student_data:qalification, timestamp=1511867458221, value=
100,alabama,madison`
AGNES,goverenment,20-10-2000,BE, column=student_data:score, timestamp=1511867458221, value=
100,alabama,madison`
AGNES,goverenment,20-10-2000,BE, column=student_data:sector, timestamp=1511867458221, value=
100,alabama,madison`
AGNES,goverenment,20-10-2000,BE, column=student_data:state, timestamp=1511867458221, value=
100,alabama,madison`
AGNEW,goverenment,20-10-2000,BC0 column=student_data:DOB, timestamp=1511867458030, value=

```